

# What is Semantic Segmentation?

Semantic segmentation refers to the process of associating each pixel in an image with a specific class label. These labels can include objects such as a person, car, flower, or a piece of furniture, among others.

We can think of semantic segmentation as a form of image classification at the pixel level. For example, in an image containing multiple cars, semantic segmentation will label all objects as 'car.' However, there's a distinct class of models called instance segmentation, which is capable of labeling individual instances of an object within an image. This type of segmentation is particularly valuable in applications that require object counting, such as monitoring foot traffic in a mall.

## Import Libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split

import torch
import torch.nn as nn
from torch.utils.data import Dataset, DataLoader
from torchvision import transforms as T
import torchvision
import torch.nn.functional as F
from torch.autograd import Variable

from PIL import Image
import cv2
import albumentations as A

import time
import os
from tqdm.notebook import tqdm

!pip install -q segmentation-models-pytorch
!pip install -q torchsummary
```

```
ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This behaviour is the source of the following dependency conflicts.
jupyterlab-git 0.11.0 requires nbdime<2.0.0,>=1.1.0, but you have nbdime 2.1.0 which is incompatible.
bokeh 2.2.3 requires tornado>=5.1, but you have tornado 5.0.2 which is incompatible.
autogluon-core 0.1.0b20210210 requires numpy==1.19, but you have numpy 1.19.5 which is incompatible.
WARNING: You are using pip version 21.0.1; however, version 23.2.1 is available.
You should consider upgrading via the '/opt/conda/bin/python3.7 -m pip install --upgrade pip' command.
WARNING: You are using pip version 21.0.1; however, version 23.2.1 is available.
You should consider upgrading via the '/opt/conda/bin/python3.7 -m pip install --upgrade pip' command.
```

## Preprocessing

```
In [2]: IMAGE_PATH = '../input/semantic-drone-dataset/dataset/semantic_drone_dataset/original'
MASK_PATH = '../input/semantic-drone-dataset/dataset/semantic_drone_dataset/label_imag
```

```
In [3]: n_classes = 23

def create_df():
    name = []
    for dirname, _, filenames in os.walk(IMAGE_PATH):
        for filename in filenames:
            name.append(filename.split('.')[0])

    return pd.DataFrame({'id': name}, index = np.arange(0, len(name)))

df = create_df()
print('Total Images: ', len(df))
```

```
Total Images: 400
```

```
In [4]: #split data
X_trainval, X_test = train_test_split(df['id'].values, test_size=0.1, random_state=19)
X_train, X_val = train_test_split(X_trainval, test_size=0.15, random_state=19)
```

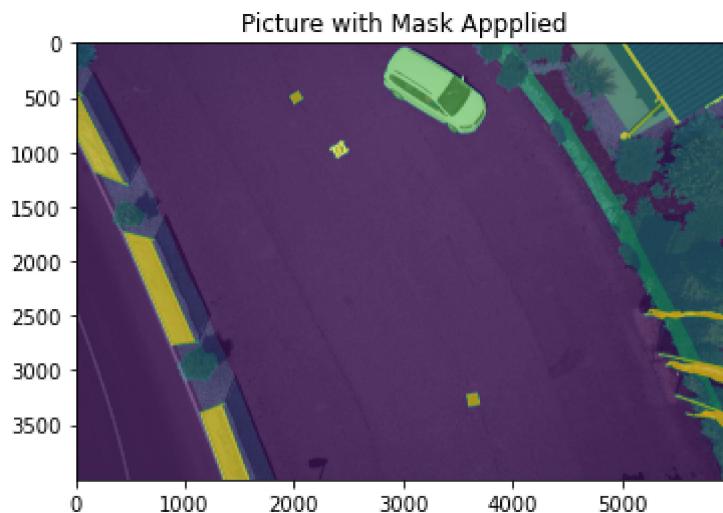
```
print('Train Size    : ', len(X_train))
print('Val Size     : ', len(X_val))
print('Test Size    : ', len(X_test))
```

```
Train Size    : 306
Val Size     : 54
Test Size    : 40
```

```
In [5]: img = Image.open(IMAGE_PATH + df['id'][100] + '.jpg')
mask = Image.open(MASK_PATH + df['id'][100] + '.png')
print('Image Size', np.asarray(img).shape)
print('Mask Size', np.asarray(mask).shape)

plt.imshow(img)
plt.imshow(mask, alpha=0.6)
plt.title('Picture with Mask Applied')
plt.show()
```

Image Size (4000, 6000, 3)  
Mask Size (4000, 6000)



## Dataset

In [6]:

```
class DroneDataset(Dataset):

    def __init__(self, img_path, mask_path, X, mean, std, transform=None, patch=False):
        self.img_path = img_path
        self.mask_path = mask_path
        self.X = X
        self.transform = transform
        self.patches = patch
        self.mean = mean
        self.std = std

    def __len__(self):
        return len(self.X)

    def __getitem__(self, idx):
        img = cv2.imread(self.img_path + self.X[idx] + '.jpg')
        img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
        mask = cv2.imread(self.mask_path + self.X[idx] + '.png', cv2.IMREAD_GRAYSCALE)

        if self.transform is not None:
            aug = self.transform(image=img, mask=mask)
            img = Image.fromarray(aug['image'])
            mask = aug['mask']

        if self.transform is None:
            img = Image.fromarray(img)

        t = T.Compose([T.ToTensor(), T.Normalize(self.mean, self.std)])
        img = t(img)
        mask = torch.from_numpy(mask).long()

        if self.patches:
            img, mask = self.tiles(img, mask)

        return img, mask
```

```

def tiles(self, img, mask):

    img_patches = img.unfold(1, 512, 512).unfold(2, 768, 768)
    img_patches = img_patches.contiguous().view(3,-1, 512, 768)
    img_patches = img_patches.permute(1,0,2,3)

    mask_patches = mask.unfold(0, 512, 512).unfold(1, 768, 768)
    mask_patches = mask_patches.contiguous().view(-1, 512, 768)

    return img_patches, mask_patches

```

```

In [7]: mean=[0.485, 0.456, 0.406]
        std=[0.229, 0.224, 0.225]

t_train = A.Compose([A.Resize(704, 1056, interpolation=cv2.INTER_NEAREST), A.HorizontalFlip(),
                     A.GridDistortion(p=0.2), A.RandomBrightnessContrast((0,0.5),(0,0.5)),
                     A.GaussNoise()])

t_val = A.Compose([A.Resize(704, 1056, interpolation=cv2.INTER_NEAREST), A.HorizontalFlip(),
                   A.GridDistortion(p=0.2)])

#datasets
train_set = DroneDataset(IMAGE_PATH, MASK_PATH, X_train, mean, std, t_train, patch=False)
val_set = DroneDataset(IMAGE_PATH, MASK_PATH, X_val, mean, std, t_val, patch=False)

#dataLoader
batch_size= 3

train_loader = DataLoader(train_set, batch_size=batch_size, shuffle=True)
val_loader = DataLoader(val_set, batch_size=batch_size, shuffle=True)

```

## Model

```

In [8]: model = smp.Unet('mobilenet_v2', encoder_weights='imagenet', classes=23, activation=None)

Downloading: "https://download.pytorch.org/models/mobilenet_v2-b0353104.pth" to /root/.cache/torch/hub/checkpoints/mobilenet_v2-b0353104.pth
0%|          | 0.00/13.6M [00:00<?, ?B/s]

```

```
In [9]: model
```

```
Out[9]: Unet(  
    (encoder): MobileNetV2Encoder(  
        (features): Sequential(  
            (0): ConvBNReLU(  
                (0): Conv2d(3, 32, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)  
            (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
            (2): ReLU6(inplace=True)  
        )  
        (1): InvertedResidual(  
            (conv): Sequential(  
                (0): ConvBNReLU(  
                    (0): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), groups=32, bias=False)  
                    (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
                    (2): ReLU6(inplace=True)  
                )  
                (1): Conv2d(32, 16, kernel_size=(1, 1), stride=(1, 1), bias=False)  
                (2): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
            )  
            (2): InvertedResidual(  
                (conv): Sequential(  
                    (0): ConvBNReLU(  
                        (0): Conv2d(16, 96, kernel_size=(1, 1), stride=(1, 1), bias=False)  
                        (1): BatchNorm2d(96, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
                    )  
                    (2): ReLU6(inplace=True)  
                )  
                (1): ConvBNReLU(  
                    (0): Conv2d(96, 96, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), groups=96, bias=False)  
                    (1): BatchNorm2d(96, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
                )  
                (2): ReLU6(inplace=True)  
            )  
            (1): Conv2d(96, 24, kernel_size=(1, 1), stride=(1, 1), bias=False)  
            (3): BatchNorm2d(24, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
        )  
        (2): InvertedResidual(  
            (conv): Sequential(  
                (0): ConvBNReLU(  
                    (0): Conv2d(24, 144, kernel_size=(1, 1), stride=(1, 1), bias=False)  
                    (1): BatchNorm2d(144, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
                )  
                (2): ReLU6(inplace=True)  
            )  
            (1): ConvBNReLU(  
                (0): Conv2d(144, 144, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), groups=144, bias=False)  
                (1): BatchNorm2d(144, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
            )  
            (2): ReLU6(inplace=True)  
        )  
        (2): Conv2d(144, 24, kernel_size=(1, 1), stride=(1, 1), bias=False)
```

```
        (3): BatchNorm2d(24, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
    )
)
(4): InvertedResidual(
    (conv): Sequential(
        (0): ConvBNReLU(
            (0): Conv2d(24, 144, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(144, eps=1e-05, momentum=0.1, affine=True, track_running
_stats=True)
            (2): ReLU6(inplace=True)
        )
        (1): ConvBNReLU(
            (0): Conv2d(144, 144, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1),
groups=144, bias=False)
            (1): BatchNorm2d(144, eps=1e-05, momentum=0.1, affine=True, track_running
_stats=True)
            (2): ReLU6(inplace=True)
        )
        (2): Conv2d(144, 32, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (3): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
    )
)
(5): InvertedResidual(
    (conv): Sequential(
        (0): ConvBNReLU(
            (0): Conv2d(32, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True, track_running
_stats=True)
            (2): ReLU6(inplace=True)
        )
        (1): ConvBNReLU(
            (0): Conv2d(192, 192, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
groups=192, bias=False)
            (1): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True, track_running
_stats=True)
            (2): ReLU6(inplace=True)
        )
        (2): Conv2d(192, 32, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (3): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
    )
)
(6): InvertedResidual(
    (conv): Sequential(
        (0): ConvBNReLU(
            (0): Conv2d(32, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True, track_running
_stats=True)
            (2): ReLU6(inplace=True)
        )
        (1): ConvBNReLU(
            (0): Conv2d(192, 192, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
groups=192, bias=False)
            (1): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True, track_running
_stats=True)
            (2): ReLU6(inplace=True)
        )
        (2): Conv2d(192, 32, kernel_size=(1, 1), stride=(1, 1), bias=False)
    )
)
```

```
        (3): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
    )
)
(7): InvertedResidual(
    (conv): Sequential(
        (0): ConvBNReLU(
            (0): Conv2d(32, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True, track_running
_stats=True)
            (2): ReLU6(inplace=True)
        )
        (1): ConvBNReLU(
            (0): Conv2d(192, 192, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1),
groups=192, bias=False)
            (1): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True, track_running
_stats=True)
            (2): ReLU6(inplace=True)
        )
        (2): Conv2d(192, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (3): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
    )
)
(8): InvertedResidual(
    (conv): Sequential(
        (0): ConvBNReLU(
            (0): Conv2d(64, 384, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(384, eps=1e-05, momentum=0.1, affine=True, track_running
_stats=True)
            (2): ReLU6(inplace=True)
        )
        (1): ConvBNReLU(
            (0): Conv2d(384, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
groups=384, bias=False)
            (1): BatchNorm2d(384, eps=1e-05, momentum=0.1, affine=True, track_running
_stats=True)
            (2): ReLU6(inplace=True)
        )
        (2): Conv2d(384, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (3): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
    )
)
(9): InvertedResidual(
    (conv): Sequential(
        (0): ConvBNReLU(
            (0): Conv2d(64, 384, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(384, eps=1e-05, momentum=0.1, affine=True, track_running
_stats=True)
            (2): ReLU6(inplace=True)
        )
        (1): ConvBNReLU(
            (0): Conv2d(384, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
groups=384, bias=False)
            (1): BatchNorm2d(384, eps=1e-05, momentum=0.1, affine=True, track_running
_stats=True)
            (2): ReLU6(inplace=True)
        )
        (2): Conv2d(384, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
    )
)
```

```
        (3): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
    )
)
(10): InvertedResidual(
    (conv): Sequential(
        (0): ConvBNReLU(
            (0): Conv2d(64, 384, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(384, eps=1e-05, momentum=0.1, affine=True, track_running
_stats=True)
            (2): ReLU6(inplace=True)
        )
        (1): ConvBNReLU(
            (0): Conv2d(384, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
groups=384, bias=False)
            (1): BatchNorm2d(384, eps=1e-05, momentum=0.1, affine=True, track_running
_stats=True)
            (2): ReLU6(inplace=True)
        )
        (2): Conv2d(384, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (3): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
    )
)
(11): InvertedResidual(
    (conv): Sequential(
        (0): ConvBNReLU(
            (0): Conv2d(64, 384, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(384, eps=1e-05, momentum=0.1, affine=True, track_running
_stats=True)
            (2): ReLU6(inplace=True)
        )
        (1): ConvBNReLU(
            (0): Conv2d(384, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
groups=384, bias=False)
            (1): BatchNorm2d(384, eps=1e-05, momentum=0.1, affine=True, track_running
_stats=True)
            (2): ReLU6(inplace=True)
        )
        (2): Conv2d(384, 96, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (3): BatchNorm2d(96, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
    )
)
(12): InvertedResidual(
    (conv): Sequential(
        (0): ConvBNReLU(
            (0): Conv2d(96, 576, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(576, eps=1e-05, momentum=0.1, affine=True, track_running
_stats=True)
            (2): ReLU6(inplace=True)
        )
        (1): ConvBNReLU(
            (0): Conv2d(576, 576, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
groups=576, bias=False)
            (1): BatchNorm2d(576, eps=1e-05, momentum=0.1, affine=True, track_running
_stats=True)
            (2): ReLU6(inplace=True)
        )
        (2): Conv2d(576, 96, kernel_size=(1, 1), stride=(1, 1), bias=False)
```

```
        (3): BatchNorm2d(96, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
    )
)
(13): InvertedResidual(
    (conv): Sequential(
        (0): ConvBNReLU(
            (0): Conv2d(96, 576, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(576, eps=1e-05, momentum=0.1, affine=True, track_running
_stats=True)
                (2): ReLU6(inplace=True)
            )
        (1): ConvBNReLU(
            (0): Conv2d(576, 576, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
groups=576, bias=False)
            (1): BatchNorm2d(576, eps=1e-05, momentum=0.1, affine=True, track_running
_stats=True)
                (2): ReLU6(inplace=True)
            )
        (2): Conv2d(576, 96, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (3): BatchNorm2d(96, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
    )
)
(14): InvertedResidual(
    (conv): Sequential(
        (0): ConvBNReLU(
            (0): Conv2d(96, 576, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(576, eps=1e-05, momentum=0.1, affine=True, track_running
_stats=True)
                (2): ReLU6(inplace=True)
            )
        (1): ConvBNReLU(
            (0): Conv2d(576, 576, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1),
groups=576, bias=False)
            (1): BatchNorm2d(576, eps=1e-05, momentum=0.1, affine=True, track_running
_stats=True)
                (2): ReLU6(inplace=True)
            )
        (2): Conv2d(576, 160, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (3): BatchNorm2d(160, eps=1e-05, momentum=0.1, affine=True, track_running_s
tats=True)
    )
)
(15): InvertedResidual(
    (conv): Sequential(
        (0): ConvBNReLU(
            (0): Conv2d(160, 960, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(960, eps=1e-05, momentum=0.1, affine=True, track_running
_stats=True)
                (2): ReLU6(inplace=True)
            )
        (1): ConvBNReLU(
            (0): Conv2d(960, 960, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
groups=960, bias=False)
            (1): BatchNorm2d(960, eps=1e-05, momentum=0.1, affine=True, track_running
_stats=True)
                (2): ReLU6(inplace=True)
            )
        (2): Conv2d(960, 160, kernel_size=(1, 1), stride=(1, 1), bias=False)
```

```

        (3): BatchNorm2d(160, eps=1e-05, momentum=0.1, affine=True, track_running_s
tats=True)
    )
)
(16): InvertedResidual(
    (conv): Sequential(
        (0): ConvBNReLU(
            (0): Conv2d(160, 960, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(960, eps=1e-05, momentum=0.1, affine=True, track_running
_stats=True)
            (2): ReLU6(inplace=True)
        )
        (1): ConvBNReLU(
            (0): Conv2d(960, 960, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
groups=960, bias=False)
            (1): BatchNorm2d(960, eps=1e-05, momentum=0.1, affine=True, track_running
_stats=True)
            (2): ReLU6(inplace=True)
        )
        (2): Conv2d(960, 160, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (3): BatchNorm2d(160, eps=1e-05, momentum=0.1, affine=True, track_running_s
tats=True)
    )
)
(17): InvertedResidual(
    (conv): Sequential(
        (0): ConvBNReLU(
            (0): Conv2d(160, 960, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(960, eps=1e-05, momentum=0.1, affine=True, track_running
_stats=True)
            (2): ReLU6(inplace=True)
        )
        (1): ConvBNReLU(
            (0): Conv2d(960, 960, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
groups=960, bias=False)
            (1): BatchNorm2d(960, eps=1e-05, momentum=0.1, affine=True, track_running
_stats=True)
            (2): ReLU6(inplace=True)
        )
        (2): Conv2d(960, 320, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (3): BatchNorm2d(320, eps=1e-05, momentum=0.1, affine=True, track_running_s
tats=True)
    )
)
(18): ConvBNReLU(
    (0): Conv2d(320, 1280, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (1): BatchNorm2d(1280, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
    (2): ReLU6(inplace=True)
)
)
)
)
decoder): UnetDecoder(
    (center): Identity()
    (blocks): ModuleList(
        (0): DecoderBlock(
            (conv1): Conv2dReLU(
                (0): Conv2d(1376, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
                (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_s

```

```
tats=True)
    (2): ReLU(inplace=True)
)
(attention1): Attention(
    (attention): Identity()
)
(conv2): Conv2dReLU(
    (0): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (2): ReLU(inplace=True)
)
(attention2): Attention(
    (attention): Identity()
)
)
(1): DecoderBlock(
    (conv1): Conv2dReLU(
        (0): Conv2d(288, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
            (2): ReLU(inplace=True)
)
(attention1): Attention(
    (attention): Identity()
)
(conv2): Conv2dReLU(
    (0): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (2): ReLU(inplace=True)
)
(attention2): Attention(
    (attention): Identity()
)
)
(2): DecoderBlock(
    (conv1): Conv2dReLU(
        (0): Conv2d(152, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
            (2): ReLU(inplace=True)
)
(attention1): Attention(
    (attention): Identity()
)
(conv2): Conv2dReLU(
    (0): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (2): ReLU(inplace=True)
)
(attention2): Attention(
    (attention): Identity()
)
```

```

        )
    (3): DecoderBlock(
        (conv1): Conv2dReLU(
            (0): Conv2d(80, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias
=False)
            (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
            (2): ReLU(inplace=True)
        )
        (attention1): Attention(
            (attention): Identity()
        )
        (conv2): Conv2dReLU(
            (0): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias
=False)
            (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
            (2): ReLU(inplace=True)
        )
        (attention2): Attention(
            (attention): Identity()
        )
    )
    (4): DecoderBlock(
        (conv1): Conv2dReLU(
            (0): Conv2d(32, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias
=False)
            (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
            (2): ReLU(inplace=True)
        )
        (attention1): Attention(
            (attention): Identity()
        )
        (conv2): Conv2dReLU(
            (0): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias
=False)
            (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, track_running_st
ats=True)
            (2): ReLU(inplace=True)
        )
        (attention2): Attention(
            (attention): Identity()
        )
    )
)
)
(segmentation_head): SegmentationHead(
    (0): Conv2d(16, 23, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): Identity()
    (2): Activation(
        (activation): Identity()
    )
)
)
)

```

# Training

```
In [10]: def pixel_accuracy(output, mask):
    with torch.no_grad():
        output = torch.argmax(F.softmax(output, dim=1), dim=1)
        correct = torch.eq(output, mask).int()
        accuracy = float(correct.sum()) / float(correct.numel())
    return accuracy
```

```
In [11]: def mIoU(pred_mask, mask, smooth=1e-10, n_classes=23):
    with torch.no_grad():
        pred_mask = F.softmax(pred_mask, dim=1)
        pred_mask = torch.argmax(pred_mask, dim=1)
        pred_mask = pred_mask.contiguous().view(-1)
        mask = mask.contiguous().view(-1)

        iou_per_class = []
        for clas in range(0, n_classes): #Loop per pixel class
            true_class = pred_mask == clas
            true_label = mask == clas

            if true_label.long().sum().item() == 0: #no exist label in this loop
                iou_per_class.append(np.nan)
            else:
                intersect = torch.logical_and(true_class, true_label).sum().float().item()
                union = torch.logical_or(true_class, true_label).sum().float().item()

                iou = (intersect + smooth) / (union +smooth)
                iou_per_class.append(iou)
    return np.nanmean(iou_per_class)
```

```
In [12]: def get_lr(optimizer):
    for param_group in optimizer.param_groups:
        return param_group['lr']

def fit(epochs, model, train_loader, val_loader, criterion, optimizer, scheduler, patch_size,
       torch.cuda.empty_cache()
       train_losses = []
       test_losses = []
       val_iou = []; val_acc = []
       train_iou = []; train_acc = []
       lrs = []
       min_loss = np.inf
       decrease = 1 ; not_improve=0

       model.to(device)
       fit_time = time.time()
       for e in range(epochs):
           since = time.time()
           running_loss = 0
           iou_score = 0
           accuracy = 0
           #training loop
           model.train()
           for i, data in enumerate(tqdm(train_loader)):
               #training phase
               image_tiles, mask_tiles = data
               if patch:
                   bs, n_tiles, c, h, w = image_tiles.size()
```

```

        image_tiles = image_tiles.view(-1,c, h, w)
        mask_tiles = mask_tiles.view(-1, h, w)

        image = image_tiles.to(device); mask = mask_tiles.to(device);
#forward
        output = model(image)
        loss = criterion(output, mask)
#evaluation metrics
        iou_score += mIoU(output, mask)
        accuracy += pixel_accuracy(output, mask)
#backward
        loss.backward()
        optimizer.step() #update weight
        optimizer.zero_grad() #reset gradient

#step the Learning rate
        lrs.append(get_lr(optimizer))
        scheduler.step()

        running_loss += loss.item()

    else:
        model.eval()
        test_loss = 0
        test_accuracy = 0
        val_iou_score = 0
#validation loop
        with torch.no_grad():
            for i, data in enumerate(tqdm(val_loader)):
                #reshape to 9 patches from single image, delete batch size
                image_tiles, mask_tiles = data

                if patch:
                    bs, n_tiles, c, h, w = image_tiles.size()

                    image_tiles = image_tiles.view(-1,c, h, w)
                    mask_tiles = mask_tiles.view(-1, h, w)

                    image = image_tiles.to(device); mask = mask_tiles.to(device);
                    output = model(image)
#evaluation metrics
                    val_iou_score += mIoU(output, mask)
                    test_accuracy += pixel_accuracy(output, mask)
#Loss
                    loss = criterion(output, mask)
                    test_loss += loss.item()

#calculatio mean for each batch
        train_losses.append(running_loss/len(train_loader))
        test_losses.append(test_loss/len(val_loader))

        if min_loss > (test_loss/len(val_loader)):
            print('Loss Decreasing.. {:.3f} >> {:.3f}'.format(min_loss, (test_loss/
            min_loss = (test_loss/len(val_loader)))
            decrease += 1
            if decrease % 5 == 0:
                print('saving model...')
                torch.save(model, 'Unet-Mobilenet_v2_mIoU-{:3f}.pt'.format(val_i

```

```

    if (test_loss/len(val_loader)) > min_loss:
        not_improve += 1
        min_loss = (test_loss/len(val_loader))
        print(f'Loss Not Decrease for {not_improve} time')
        if not_improve == 7:
            print('Loss not decrease for 7 times, Stop Training')
            break

    #iou
    val_iou.append(val_iou_score/len(val_loader))
    train_iou.append(iou_score/len(train_loader))
    train_acc.append(accuracy/len(train_loader))
    val_acc.append(test_accuracy/ len(val_loader))
    print("Epoch:{}{}..".format(e+1, epochs),
          "Train Loss: {:.3f}..".format(running_loss/len(train_loader)),
          "Val Loss: {:.3f}..".format(test_loss/len(val_loader)),
          "Train mIoU:{:.3f}..".format(iou_score/len(train_loader)),
          "Val mIoU: {:.3f}..".format(val_iou_score/len(val_loader)),
          "Train Acc:{:.3f}..".format(accuracy/len(train_loader)),
          "Val Acc:{:.3f}..".format(test_accuracy/len(val_loader)),
          "Time: {:.2f}m".format((time.time()-since)/60))

    history = {'train_loss' : train_losses, 'val_loss': test_losses,
               'train_miou' :train_iou, 'val_miou':val_iou,
               'train_acc' :train_acc, 'val_acc':val_acc,
               'lrs': lrs}
    print('Total time: {:.2f} m'.format((time.time()- fit_time)/60))
    return history

```

```

In [13]: max_lr = 1e-3
epoch = 15
weight_decay = 1e-4

criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.AdamW(model.parameters(), lr=max_lr, weight_decay=weight_decay)
sched = torch.optim.lr_scheduler.OneCycleLR(optimizer, max_lr, epochs=epoch,
                                             steps_per_epoch=len(train_loader))

history = fit(epoch, model, train_loader, val_loader, criterion, optimizer, sched)

0% | 0/102 [00:00<?, ?it/s]
0% | 0/18 [00:00<?, ?it/s]
Loss Decreasing.. inf >> 2.201
Epoch:1/15.. Train Loss: 2.891.. Val Loss: 2.201.. Train mIoU:0.051.. Val mIoU: 0.11
9.. Train Acc:0.189.. Val Acc:0.500.. Time: 4.62m
0% | 0/102 [00:00<?, ?it/s]
0% | 0/18 [00:00<?, ?it/s]
Loss Decreasing.. 2.201 >> 1.497
Epoch:2/15.. Train Loss: 1.974.. Val Loss: 1.497.. Train mIoU:0.118.. Val mIoU: 0.15
2.. Train Acc:0.556.. Val Acc:0.645.. Time: 4.35m
0% | 0/102 [00:00<?, ?it/s]
0% | 0/18 [00:00<?, ?it/s]
Loss Decreasing.. 1.497 >> 1.198
Epoch:3/15.. Train Loss: 1.473.. Val Loss: 1.198.. Train mIoU:0.118.. Val mIoU: 0.14
2.. Train Acc:0.578.. Val Acc:0.642.. Time: 4.36m
0% | 0/102 [00:00<?, ?it/s]
0% | 0/18 [00:00<?, ?it/s]

```

Loss Decreasing.. 1.198 >> 1.016  
saving model...  
Epoch:4/15.. Train Loss: 1.337.. Val Loss: 1.016.. Train mIoU:0.144.. Val mIoU: 0.17  
8.. Train Acc:0.607.. Val Acc:0.700.. Time: 4.38m  
  0% | 0/102 [00:00<?, ?it/s]  
  0% | 0/18 [00:00<?, ?it/s]  
Loss Decreasing.. 1.016 >> 0.997  
Epoch:5/15.. Train Loss: 1.273.. Val Loss: 0.997.. Train mIoU:0.155.. Val mIoU: 0.19  
0.. Train Acc:0.617.. Val Acc:0.716.. Time: 4.38m  
  0% | 0/102 [00:00<?, ?it/s]  
  0% | 0/18 [00:00<?, ?it/s]  
Loss Decreasing.. 0.997 >> 0.947  
Epoch:6/15.. Train Loss: 1.125.. Val Loss: 0.947.. Train mIoU:0.179.. Val mIoU: 0.20  
8.. Train Acc:0.666.. Val Acc:0.710.. Time: 4.38m  
  0% | 0/102 [00:00<?, ?it/s]  
  0% | 0/18 [00:00<?, ?it/s]  
Loss Decreasing.. 0.947 >> 0.782  
Epoch:7/15.. Train Loss: 1.029.. Val Loss: 0.782.. Train mIoU:0.211.. Val mIoU: 0.23  
8.. Train Acc:0.701.. Val Acc:0.768.. Time: 4.42m  
  0% | 0/102 [00:00<?, ?it/s]  
  0% | 0/18 [00:00<?, ?it/s]  
Loss Not Decrease for 1 time  
Epoch:8/15.. Train Loss: 0.977.. Val Loss: 0.814.. Train mIoU:0.219.. Val mIoU: 0.21  
7.. Train Acc:0.711.. Val Acc:0.764.. Time: 4.36m  
  0% | 0/102 [00:00<?, ?it/s]  
  0% | 0/18 [00:00<?, ?it/s]  
Loss Decreasing.. 0.814 >> 0.742  
Epoch:9/15.. Train Loss: 0.916.. Val Loss: 0.742.. Train mIoU:0.233.. Val mIoU: 0.25  
7.. Train Acc:0.732.. Val Acc:0.774.. Time: 4.42m  
  0% | 0/102 [00:00<?, ?it/s]  
  0% | 0/18 [00:00<?, ?it/s]  
Loss Decreasing.. 0.742 >> 0.691  
saving model...  
Epoch:10/15.. Train Loss: 0.808.. Val Loss: 0.691.. Train mIoU:0.259.. Val mIoU: 0.27  
1.. Train Acc:0.763.. Val Acc:0.788.. Time: 4.36m  
  0% | 0/102 [00:00<?, ?it/s]  
  0% | 0/18 [00:00<?, ?it/s]  
Loss Decreasing.. 0.691 >> 0.635  
Epoch:11/15.. Train Loss: 0.749.. Val Loss: 0.635.. Train mIoU:0.274.. Val mIoU: 0.29  
2.. Train Acc:0.780.. Val Acc:0.810.. Time: 4.37m  
  0% | 0/102 [00:00<?, ?it/s]  
  0% | 0/18 [00:00<?, ?it/s]  
Loss Not Decrease for 2 time  
Epoch:12/15.. Train Loss: 0.699.. Val Loss: 0.636.. Train mIoU:0.295.. Val mIoU: 0.30  
0.. Train Acc:0.796.. Val Acc:0.806.. Time: 4.39m  
  0% | 0/102 [00:00<?, ?it/s]  
  0% | 0/18 [00:00<?, ?it/s]  
Loss Decreasing.. 0.636 >> 0.602  
Epoch:13/15.. Train Loss: 0.678.. Val Loss: 0.602.. Train mIoU:0.298.. Val mIoU: 0.31  
7.. Train Acc:0.801.. Val Acc:0.821.. Time: 4.41m  
  0% | 0/102 [00:00<?, ?it/s]  
  0% | 0/18 [00:00<?, ?it/s]  
Loss Decreasing.. 0.602 >> 0.579  
Epoch:14/15.. Train Loss: 0.634.. Val Loss: 0.579.. Train mIoU:0.318.. Val mIoU: 0.32  
1.. Train Acc:0.814.. Val Acc:0.825.. Time: 4.43m  
  0% | 0/102 [00:00<?, ?it/s]  
  0% | 0/18 [00:00<?, ?it/s]

```
Loss Decreasing.. 0.579 >> 0.574
Epoch:15/15.. Train Loss: 0.650.. Val Loss: 0.574.. Train mIoU:0.314.. Val mIoU: 0.31
7.. Train Acc:0.812.. Val Acc:0.825.. Time: 4.39m
Total time: 66.03 m
```

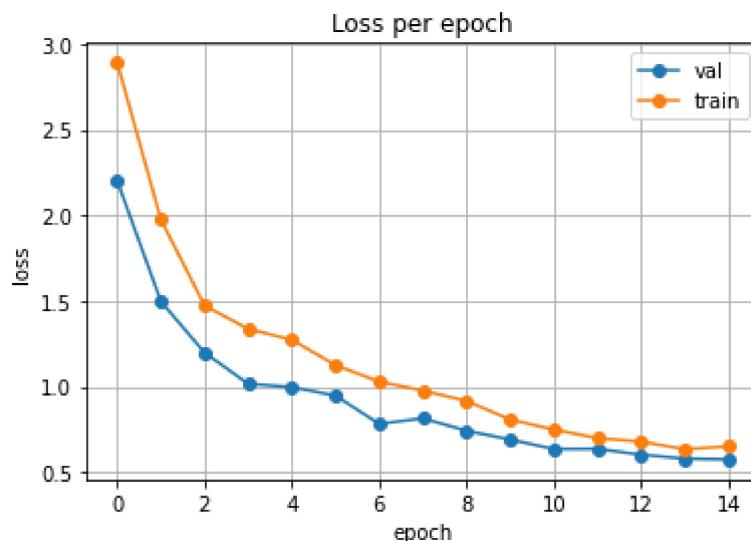
```
In [14]: torch.save(model, 'Unet-Mobilenet.pt')
```

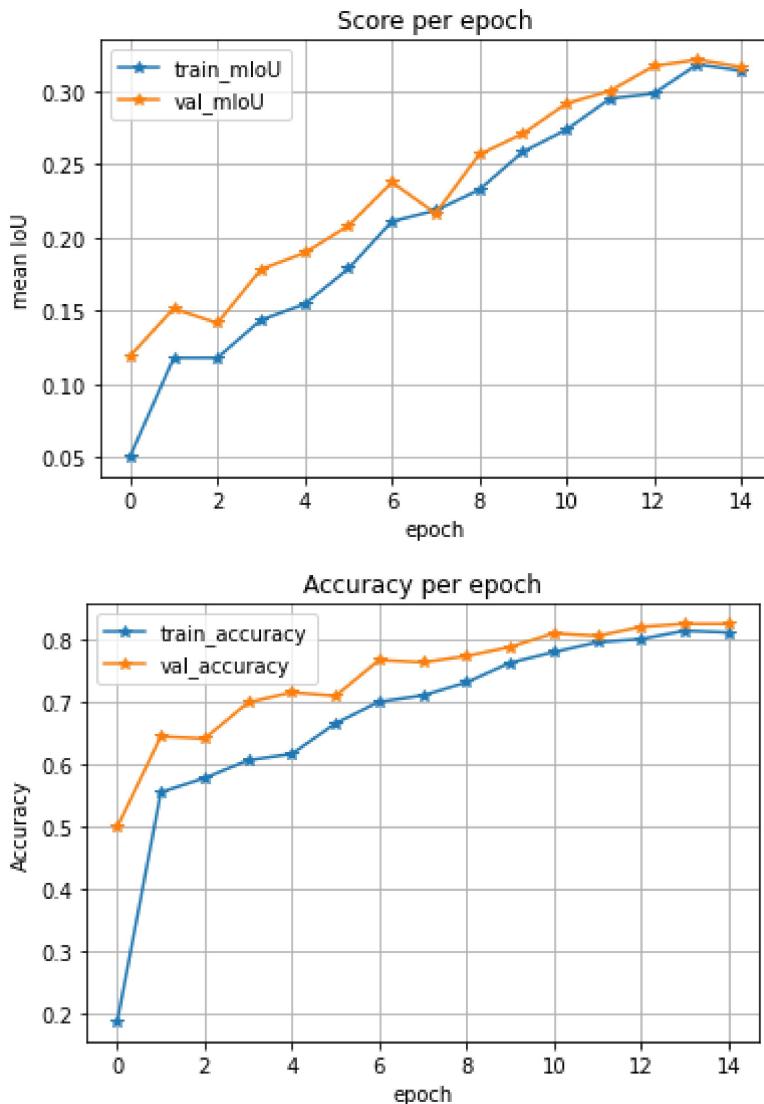
```
def plot_loss(history):
    plt.plot(history['val_loss'], label='val', marker='o')
    plt.plot( history['train_loss'], label='train', marker='o')
    plt.title('Loss per epoch'); plt.ylabel('loss');
    plt.xlabel('epoch')
    plt.legend(), plt.grid()
    plt.show()

def plot_score(history):
    plt.plot(history['train_miou'], label='train_mIoU', marker='*')
    plt.plot(history['val_miou'], label='val_mIoU', marker='*')
    plt.title('Score per epoch'); plt.ylabel('mean IoU')
    plt.xlabel('epoch')
    plt.legend(), plt.grid()
    plt.show()

def plot_acc(history):
    plt.plot(history['train_acc'], label='train_accuracy', marker='*')
    plt.plot(history['val_acc'], label='val_accuracy', marker='*')
    plt.title('Accuracy per epoch'); plt.ylabel('Accuracy')
    plt.xlabel('epoch')
    plt.legend(), plt.grid()
    plt.show()
```

```
In [16]: plot_loss(history)
plot_score(history)
plot_acc(history)
```





## Evaluation

```
In [17]: class DroneTestDataset(Dataset):

    def __init__(self, img_path, mask_path, X, transform=None):
        self.img_path = img_path
        self.mask_path = mask_path
        self.X = X
        self.transform = transform

    def __len__(self):
        return len(self.X)

    def __getitem__(self, idx):
        img = cv2.imread(self.img_path + self.X[idx] + '.jpg')
        img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
        mask = cv2.imread(self.mask_path + self.X[idx] + '.png', cv2.IMREAD_GRAYSCALE)

        if self.transform is not None:
            aug = self.transform(image=img, mask=mask)
            img = Image.fromarray(aug['image'])
            mask = aug['mask']

        return {'img': img, 'mask': mask}
```

```

        if self.transform is None:
            img = Image.fromarray(img)

        mask = torch.from_numpy(mask).long()

    return img, mask

t_test = A.Resize(768, 1152, interpolation=cv2.INTER_NEAREST)
test_set = DroneTestDataset(IMAGE_PATH, MASK_PATH, X_test, transform=t_test)

```

## Result

```
In [18]: def predict_image_mask_miou(model, image, mask, mean=[0.485, 0.456, 0.406], std=[0.229,
    model.eval()
    t = T.Compose([T.ToTensor(), T.Normalize(mean, std)])
    image = t(image)
    model.to(device); image=image.to(device)
    mask = mask.to(device)
    with torch.no_grad():

        image = image.unsqueeze(0)
        mask = mask.unsqueeze(0)

        output = model(image)
        score = mIoU(output, mask)
        masked = torch.argmax(output, dim=1)
        masked = masked.cpu().squeeze(0)
    return masked, score
```

```
In [19]: def predict_image_mask_pixel(model, image, mask, mean=[0.485, 0.456, 0.406], std=[0.229,
    model.eval()
    t = T.Compose([T.ToTensor(), T.Normalize(mean, std)])
    image = t(image)
    model.to(device); image=image.to(device)
    mask = mask.to(device)
    with torch.no_grad():

        image = image.unsqueeze(0)
        mask = mask.unsqueeze(0)

        output = model(image)
        acc = pixel_accuracy(output, mask)
        masked = torch.argmax(output, dim=1)
        masked = masked.cpu().squeeze(0)
    return masked, acc
```

```
In [20]: image, mask = test_set[3]
pred_mask, score = predict_image_mask_miou(model, image, mask)
```

```
In [21]: def miou_score(model, test_set):
    score_iou = []
    for i in tqdm(range(len(test_set))):
        img, mask = test_set[i]
        pred_mask, score = predict_image_mask_miou(model, img, mask)
```

```
        score_iou.append(score)
    return score_iou
```

```
In [22]: mob_miou = miou_score(model, test_set)
```

```
0% | 0/40 [00:00<?, ?it/s]
```

```
In [23]: def pixel_acc(model, test_set):
    accuracy = []
    for i in tqdm(range(len(test_set))):
        img, mask = test_set[i]
        pred_mask, acc = predict_image_mask_pixel(model, img, mask)
        accuracy.append(acc)
    return accuracy
```

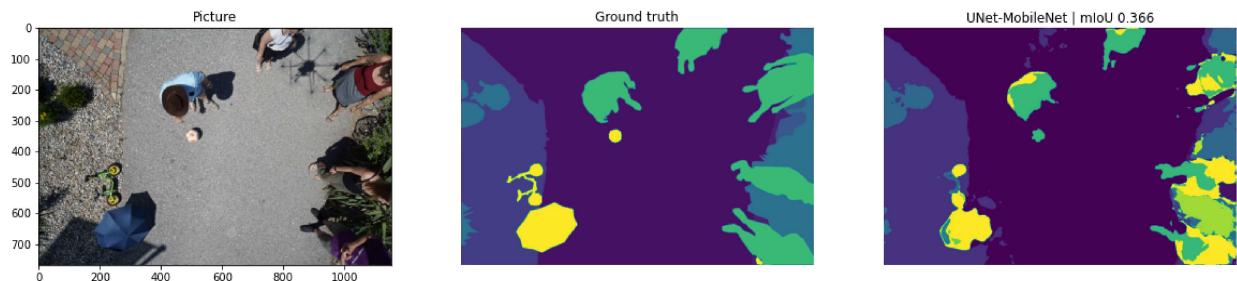
```
In [24]: mob_acc = pixel_acc(model, test_set)
```

```
0% | 0/40 [00:00<?, ?it/s]
```

```
In [25]: fig, (ax1, ax2, ax3) = plt.subplots(1,3, figsize=(20,10))
ax1.imshow(image)
ax1.set_title('Picture');

ax2.imshow(mask)
ax2.set_title('Ground truth')
ax2.set_axis_off()

ax3.imshow(pred_mask)
ax3.set_title('UNet-MobileNet | mIoU {:.3f}'.format(score))
ax3.set_axis_off()
```



```
In [26]: image2, mask2 = test_set[4]
pred_mask2, score2 = predict_image_mask_miou(model, image2, mask2)

fig, (ax1, ax2, ax3) = plt.subplots(1,3, figsize=(20,10))
ax1.imshow(image2)
ax1.set_title('Picture');

ax2.imshow(mask2)
ax2.set_title('Ground truth')
ax2.set_axis_off()

ax3.imshow(pred_mask2)
ax3.set_title('UNet-MobileNet | mIoU {:.3f}'.format(score2))
ax3.set_axis_off()
```

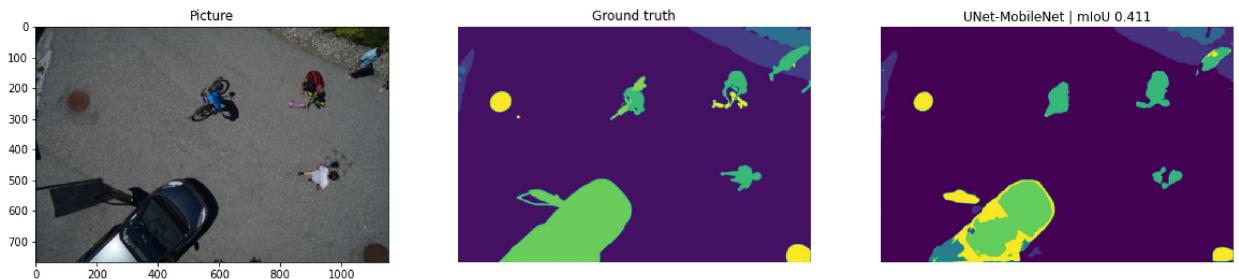


```
In [27]: image3, mask3 = test_set[6]
pred_mask3, score3 = predict_image_mask_miou(model, image3, mask3)

fig, (ax1, ax2, ax3) = plt.subplots(1,3, figsize=(20,10))
ax1.imshow(image3)
ax1.set_title('Picture');

ax2.imshow(mask3)
ax2.set_title('Ground truth')
ax2.set_axis_off()

ax3.imshow(pred_mask3)
ax3.set_title('UNet-MobileNet | mIoU {:.3f}'.format(score3))
ax3.set_axis_off()
```



```
In [28]: print('Test Set mIoU', np.mean(mob_miou))
```

Test Set mIoU 0.3526479529490619

```
In [29]: print('Test Set Pixel Accuracy', np.mean(mob_acc))
```

Test Set Pixel Accuracy 0.8131760208695024